

What and How Does the Translation Task Work Inside an LLM?

An interpretability study across eight multilingual models, and what it implies for quantization

Arctos project — Phase One findings report

2026-05-28

Abstract

We investigate *how* eight open multilingual language models carry out machine translation, using five complementary interpretability methods on a single uniform hooking framework, and a sixth purpose-built “language-pivot” analysis aimed squarely at the title question. Across decoder-only models spanning four lineages, two normalizations, two positional encodings, and 2022–2025 generations — plus one encoder–decoder (NLLB) — a consistent mechanistic picture emerges: translation is computed as a depth-ordered pipeline (source encoding $\rightarrow$ a language-agnostic / pivot representation in the middle $\rightarrow$ target-language commitment in the final layers), and the *relative depth* of this pipeline generalizes across architectures with a single clean exception (Gemma-family). We also report a negative result that matters for compression: component *importance* (where the MT work concentrates) does *not* predict quantization *sensitivity* (where numerical precision matters), confirmed across two independent metrics. We close with what this means for an interpretability-guided quantization method.

1 The question, and how we attack it

The title question — *what and how does translation work in an LLM?* — is mechanistic: not “how well does it translate” but “what internal computation produces the translation, and where in the network does each part happen.” We answer it by reading and intervening on the **residual stream** [7], the additive running sum that flows down a transformer [21]:

$$\text{resid}_\ell = \text{embed} + \sum_{\ell' \leq \ell} (\text{attn_out}_{\ell'} + \text{mlp_out}_{\ell'}).$$

Every method below either *reads* this stream (logit lens, probing, pivot trajectory), *decomposes* it into per-component contributions (IFR, DLA), or *intervenes* on it causally (attribution patching, noise injection). A single wrapper (`HookedModel`) exposes a uniform interface — `run_with_cache`, `W_U`, `W_O`, `ln_final` — across all architectures, so the same method code runs on Cohere, Llama, Gemma, BLOOM, and (with a dedicated path) the NLLB encoder–decoder.

2 Setup

Models. Eight models chosen to span training intent, architecture family, and generation (Table 1). The TowerBase $\rightarrow$ TowerInstruct pair is a controlled ablation isolating MT-task supervised fine-tuning (SFT); both build on a continued-pretraining (CPT) base.

Model	Year	Architecture	Training intent
Aya Expanse 8B [5]	2024	Cohere, RoPE, RMSNorm	general multilingual LM
TowerBase 7B [2]	2024	Llama-2	+ bilingual CPT (no SFT)
TowerInstruct 7B [2]	2024	Llama-2	+ MT-task SFT
Tower-Plus 9B [20]	2025	Gemma 2 [10]	CPT + MT-SFT
BLOOM 7B1 [4]	2022	ALiBi, LayerNorm	multilingual pretraining
EuroLLM 9B [16]	2024	Llama-3	European MT specialist
Llama-3.1 8B [12]	2024	Llama-3, GQA	general LM
Gemma-3 12B [11]	2025	Gemma 3	baseline (Google QAT)
NLLB-200 3.3B [17]	2022	encoder-decoder	MT-purpose-built

Table 1: The model set. Bold marks the architectural outliers that test generalization.

Data. FLORES+ dev [17], three language pairs chosen for contrast: **cs**→**de** (same script, sanity pair), **en**→**zh** (cross-script, Han target), **en**→**arz** (cross-script, Egyptian Arabic — a low-resource, non-standard variety). 200 sentences/pair for the correlational methods; fewer for the generation-heavy causal ones.

3 The methods, and what each reveals

Logit lens. Apply the model’s own final norm + unembedding to an *intermediate* residual: $\text{logits}_\ell = \text{LN}_f(\text{resid}_\ell) W_U$ [18]. This reads “what would the model predict if it had to decide at layer ℓ .” Tracking the probability mass on the gold target tokens across layers shows *when* the model commits to the translation. (Invariant we verify: at the last layer the lens equals the model’s real logits, to $< 10^{-4}$.) A learned refinement, the tuned lens [3], trades calibration for added trainable parameters; we use the parameter-free logit lens throughout.

Probing. Train a linear classifier per layer [1] to predict a feature (source/target language) from the residual; report *selectivity* = accuracy – control-task accuracy [13], which separates “the feature is represented” from “the probe memorized.” Reveals *where* language identity becomes linearly decodable.

Information Flow Routes (IFR). Following Ferrando and Voita [8], decompose the residual into per-component contributions and take their L_1 -normalized magnitudes: $\text{IFR}(c) = \|c\|_1 / \sum_{c'} \|c'\|_1$, averaged over the set. Unsigned; tells you which components are *loud*. Reveals *where* (which layers/heads) MT computation concentrates.

Direct logit attribution (DLA). Decompose the *target-token logit* into signed per-component contributions via the norm-linearized unembed [7]: $\text{DLA}(c) = (c \odot s) \cdot W_U[:, t]$. Signed: positive = pushes toward the target token (a “predictor”), negative = away (a “suppressor”). Reveals *which* components are causally aimed at the right output token.

Attribution patching (causal). First-order approximation of activation patching [14]: $\text{effect}(c) \approx \partial L / \partial a_c \cdot (a_c^{\text{corrupt}} - a_c^{\text{clean}})$. One backward pass scores every component. Reveals which components, if damaged, *break* the translation.

Model	source-ID peak layer (sel.)	IFR 1st-quarter	IFR last-quarter
Aya Expanse 8B	L17 (0.55)	8%	56%
TowerBase 7B	L0 (0.60)	6%	45%
TowerInstruct 7B	L0 (0.56)	6%	46%
BLOOM 7B1	L0 (0.55)	12%	44%
EuroLLM 9B	L9 (0.54)	2%	52%
Llama-3.1 8B	L0 (0.55)	8%	47%
NLLB-200 3.3B (decoder)	—	12%	45%
Tower-Plus 9B (Gemma 2)	L7 (0.54)	32%	17%
Gemma-3 12B	L29 (0.59)	92%	3%

Table 2: Q1/Q4 depth structure. Six decoder-only models *and* the NLLB encoder–decoder share the profile (first quarter light, last quarter dominant). Both Gemma-family models invert it — an architecture artifact of Gemma’s extra per-layer norms + logit soft-capping.

Noise sensitivity (quantization proxy). Add relative Gaussian noise to a component’s weights and measure the resulting quality drop — directly simulating quantization error [9, 15]. Reveals where numerical *precision* matters.

Language-pivot trajectory (the “how” centerpiece). For cross-script pairs the target uses a different writing system from English/source (Han for zh, Arabic for arz, vs. Latin). We logit-lens each layer and classify the script of the top-50 tokens, measuring the share of probability mass on Latin vs. target script vs. other. If the model “thinks in English/source” mid-stack and converts to the target script late, the trajectory shows Latin dominance giving way to target script in the final layers — the mechanistic signature of pivot-based translation [cf. 6, 22].

4 Findings: what and how

4.1 Translation is a depth-ordered pipeline

Three methods agree on the coarse structure. **(1) Source language** is linearly decodable from very early — typically the embedding (L0) for heavily multilingually-trained models and mid-network (L9–L17) for general-LM-style models (Table 2), but present at 0.4–0.55 selectivity at *every* layer. **(2) Target commitment** (logit-lens mass on the gold target tokens) stays near zero through the first half and rises only in the final layers (Fig. 1). **(3) Computation magnitude** (IFR) concentrates in the last quarter of depth (Fig. 2). The picture: *understand the source early, do the cross-lingual work in the middle, commit to the target late*.

4.2 How the cross-lingual step works: the pivot trajectory

This is the most direct evidence for *how* translation happens. For the cross-script pairs we measure, per layer, the share of the top-50 lens tokens’ probability mass on the target script (Han / Arabic) versus Latin (English/source). The result is unambiguous (Table 3, Fig. 3): **target-script mass is essentially zero through the first ~80–95% of depth, then rises sharply in the final layers**. The model represents the in-progress translation in a Latin/pivot space — not the target script — for almost the entire forward pass, and converts to the target writing system only at the very end.

The en→arz pair tells the same story where the model can produce Arabic at all (crossover at depth 0.81–0.97 for the capable models; weaker models never cross, consistent with their low en→arz quality). This directly answers the “how”: *the middle of the network is not working in the*




Figure 1: Logit-lens target-token mass vs. depth fraction, all models, all pairs. Mass is near zero through mid-depth and rises late — target-language commitment is a late-layer event.




Figure 2: IFR contribution share by depth quarter. Last-quarter dominance is consistent across lineages, normalizations, positional encodings, and the decoder-only/enc-dec divide — the lone exceptions are the two Gemma-family models.

target language at all — it operates in a script-agnostic/pivot representation, and target-language identity is written in only at the end, exactly where the logit-lens target mass (Fig. 1), positive DLA (Fig. 4), and IFR magnitude (Fig. 2) all concentrate.

4.3 Which components, and what they do (DLA + attention)

DLA isolates recurring, pair-agnostic heads. In Aya, head L30.H8 is a strong cross-pair *predictor* (DLA +0.72 on cs→de, +0.44 on en→arz); L29.H16 and L26.H5 are systematic *suppressors* (negative across pairs). The signed DLA depth curve (Fig. 4) shows the same late-commitment shape as the lens, now direction-aware: positive contribution to the target token accumulates in the last layers. Attention-pattern inspection of the top-DLA heads characterizes what they do (source-attender / previous-token / sink); per-model role tables are in `results/{model}/q2/attention/head.roles.json`.

4.4 The MT footprint generalizes across architectures (Q4)

The depth signature (Table 2, Figs. 1–2) holds across: two normalizations (RMSNorm, LayerNorm), two positional encodings (RoPE, ALiBi), four lineages (Cohere, Llama, BigScience, EuroLLM), 2022→2024 generations, *and* the decoder-only ↔ encoder-decoder divide (NLLB’s decoder matches). This supports the strong form of the shared-depth hypothesis (a characteristic relative-depth signature), with one boundary: **Gemma-family inverts it**, so the signature is family-bounded, not universal.

Model (en→zh)	tgt-script @ L0	@ mid-depth	@ last	Latin→tgt crossover
Aya Expanse 8B	0.00	0.44	0.87	depth 0.50
EuroLLM 9B	0.00	0.00	0.83	depth 0.60
Tower-Plus 9B	0.00	0.00	0.56	depth 0.95
Llama-3.1 8B	0.00	0.00	0.41	depth 0.97
Gemma-3 12B	0.00	0.00	0.72	depth 0.96
BLOOM 7B1	0.00	0.00	0.61	never > Latin

Table 3: Language-pivot trajectory, en→zh. Target-script (Han) probability mass under the logit lens is near zero until late depth; the crossover where target script overtakes Latin happens in the final layers (depth fraction ≥ 0.9 for most models). Aya, the strongest general multilingual model, commits earliest (depth 0.50). This is the signature of “process in a pivot/English-aligned space, emit in the target language last.”




Figure 3: Pivot trajectories, en→zh. Left: Aya (commits to Han mid-stack). Right: Llama-3.1 (Latin/pivot dominant until the final layers, then a sharp switch to Han). In both, the target script is absent for most of the network — the cross-lingual conversion is a late, localized event.

4.5 Importance does not predict quantization sensitivity (Q5)

The components IFR/DLA flag as important are *not* the ones where quantization precision matters. Tested at two levels, both null (Table 4):

5 Synthesis: what and how translation works

Putting it together, translation inside these LLMs is a **depth-ordered pipeline**:

1. **Source encoding (early layers).** Source-language identity is linearly present from the embedding onward; early layers build a representation of the source content.
2. **Cross-lingual / pivot processing (middle layers).** The model operates in a representation not yet committed to the target script — Section 4.2 shows target-script mass is low here while Latin/pivot mass dominates. Little target-token probability has formed yet.
3. **Target commitment (final layers).** Logit-lens target mass, DLA positive contribution, and IFR magnitude all spike in the last quarter; the script flips to the target. This is where “which target word” is decided.

The *relative depths* of these stages are remarkably stable across very different architectures (Q4), suggesting the pipeline is a property of the MT task as learned by transformers, not of any one

Reproduction

All methods are in `src/interp/`; per-question runners and SLURM scripts in `experiments/`; the eight runnable tutorials (theory + mechanics) in `notebooks/`; per-question findings in `docs/findings/`.

References

- [1] Guillaume Alain and Yoshua Bengio. Understanding intermediate layers using linear classifier probes. *arXiv preprint arXiv:1610.01644*, 2016.
- [2] Duarte M. Alves, José Pombal, Nuno M. Guerreiro, Pedro H. Martins, João Alves, Amin Farajian, Ben Peters, Ricardo Rei, Patrick Fernandes, Sweta Agrawal, Pierre Colombo, José G. C. de Souza, and André F. T. Martins. Tower: An open multilingual large language model for translation-related tasks. In *Conference on Language Modeling (COLM)*, 2024. arXiv:2402.17733.
- [3] Nora Belrose, Zach Furman, Logan Smith, Danny Halawi, Igor Ostrovsky, Lev McKinney, Stella Biderman, and Jacob Steinhardt. Eliciting latent predictions from transformers with the tuned lens. *arXiv preprint arXiv:2303.08112*, 2023.
- [4] BigScience Workshop, Teven Le Scao, Angela Fan, Christopher Akiki, et al. BLOOM: A 176B-parameter open-access multilingual language model. *arXiv preprint arXiv:2211.05100*, 2022.
- [5] John Dang, Shivalika Singh, Daniel D’souza, Arash Ahmadian, Alejandro Salamanca, Madeline Smith, Aidan Peppin, Sungjin Hong, Sumanth Govindassamy, Terrence Zhao, et al. Aya expande: Combining research breakthroughs for a new multilingual frontier. *arXiv preprint arXiv:2412.04261*, 2024.
- [6] Clément Dumas, Chris Wendler, Veniamin Veselovsky, Giovanni Monea, and Robert West. Separating tongue from thought: Activation patching reveals language-agnostic concept representations in transformers. *arXiv preprint arXiv:2411.08745*, 2024.
- [7] Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, et al. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. <https://transformer-circuits.pub/2021/framework/index.html>.
- [8] Javier Ferrando and Elena Voita. Information flow routes: Automatically interpreting language models at scale. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024. arXiv:2403.00824.
- [9] Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2210.17323.
- [10] Gemma Team. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*, 2024.
- [11] Gemma Team. Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*, 2025.
- [12] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

- [13] John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP-IJCNLP)*, pages 2733–2743, 2019. doi: 10.18653/v1/D19-1275.
- [14] János Kramár, Tom Lieberum, Rohin Shah, and Neel Nanda. AtP*: An efficient and scalable method for localizing LLM behaviour to components. *arXiv preprint arXiv:2403.00745*, 2024.
- [15] Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. AWQ: Activation-aware weight quantization for LLM compression and acceleration. In *Proceedings of Machine Learning and Systems (MLSys)*, 2024. arXiv:2306.00978.
- [16] Pedro Henrique Martins, Patrick Fernandes, João Alves, Nuno M. Guerreiro, Ricardo Rei, Duarte M. Alves, José Pombal, Amin Farajian, Manuel Faysse, Mateusz Klimaszewski, Pierre Colombo, Barry Haddow, José G. C. de Souza, Alexandra Birch, and André F. T. Martins. EuroLLM: Multilingual language models for europe. *arXiv preprint arXiv:2409.16235*, 2024.
- [17] NLLB Team, Marta R. Costa-jussà, James Cross, Onur Çelebi, Maha Elbayad, Kenneth Heafield, Kevin Heffernan, et al. No language left behind: Scaling human-centered machine translation. *arXiv preprint arXiv:2207.04672*, 2022.
- [18] nostalgebraist. Interpreting GPT: The logit lens. LessWrong, 2020. <https://www.lesswrong.com/posts/AcKRB8wDpdaN6v6ru/interpreting-gpt-the-logit-lens>.
- [19] Maja Popović. chrF++: Words helping character n-grams. In *Proceedings of the Second Conference on Machine Translation (WMT)*, pages 612–618, 2017. doi: 10.18653/v1/W17-4770.
- [20] Ricardo Rei, Nuno M. Guerreiro, José Pombal, Daan van Stigt, Marcos Treviso, Luísa Coheur, José G. C. de Souza, and André F. T. Martins. Tower+: Bridging generality and translation specialization in multilingual LLMs. *arXiv preprint arXiv:2506.17080*, 2025.
- [21] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017. arXiv:1706.03762.
- [22] Chris Wendler, Veniamin Veselovsky, Giovanni Monea, and Robert West. Do llamas work in English? on the latent language of multilingual transformers. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 15366–15394, 2024. arXiv:2402.10588.